

Introduction to Audio Content Analysis

Module 3.7: Feature Postprocessing

alexander lerch

introduction

overview

corresponding textbook section

Section 3.7

■ lecture content

- derived features
- feature aggregation
- feature normalization
- problems of dimensionality
- feature selection
- feature transformation/mapping

■ learning objectives

- discuss the advantages of specific derived features
- summarize the principles of feature aggregation
- list two forms of feature normalization and explain their usefulness
- describe potential challenges with high-dimensional feature spaces



introduction

overview

corresponding textbook section

Section 3.7

■ lecture content

- derived features
- feature aggregation
- feature normalization
- problems of dimensionality
- feature selection
- feature transformation/mapping

■ learning objectives

- discuss the advantages of specific derived features
- summarize the principles of feature aggregation
- list two forms of feature normalization and explain their usefulness
- describe potential challenges with high-dimensional feature spaces



feature post-processing

introduction 1/2

- extracting multiple instantaneous features leads to
 - one feature vector per block, or
 - one feature matrix per audio file

$$\begin{aligned} \mathbf{V} &= [\mathbf{v}(0) \ \mathbf{v}(1) \ \dots \ \mathbf{v}(\mathcal{N}-1)] \\ &= \begin{bmatrix} v_0(0) & v_0(1) & \dots & v_0(\mathcal{N}-1) \\ v_1(0) & v_1(1) & \dots & v_1(\mathcal{N}-1) \\ \vdots & \vdots & \ddots & \vdots \\ v_{\mathcal{F}-1}(0) & v_{\mathcal{F}-1}(1) & \dots & v_{\mathcal{F}-1}(\mathcal{N}-1) \end{bmatrix} \end{aligned}$$

dimensions: $\mathcal{F} \times \mathcal{N}$ (number of features and number of blocks, resp.)

feature post-processing

introduction 2/2

multiple options for feature matrix processing:

- 1 derive additional features
- 2 aggregate existing features (e.g., one feature vector per file)
- 3 ensure similar scale and distribution

feature post-processing

examples of derived features

- **diff**: use the change in value

$$v_{j,\Delta}(n) = v_j(n) - v_j(n-1)$$

- **smoothed**: remove high frequency content by low-pass filtering
 - (anticausal) single-pole

$$v_{j,\text{LP}}(n) = (1 - \alpha) \cdot v_j(n) - \alpha \cdot v_{j,\text{LP}}(n-1)$$

- moving average

feature post-processing

examples of derived features

- **diff**: use the change in value

$$v_{j,\Delta}(n) = v_j(n) - v_j(n-1)$$

- **smoothed**: remove high frequency content by low-pass filtering
 - (anticausal) single-pole

$$v_{j,\text{LP}}(n) = (1 - \alpha) \cdot v_j(n) - \alpha \cdot v_{j,\text{LP}}(n-1)$$

- moving average

feature post-processing

feature normalization

■ reasons

- features have different ranges and distributions
- ensure that one feature does not have outsized impact

■ z-score normalization

$$v_{j,N}(n) = \frac{v_j(n) - \mu_{v_j}}{\sigma_{v_j}}.$$

■ min-max normalization

$$v_{j,N}(n) = \frac{v_j(n) - \min(v_j)}{\max(v_j) - \min(v_j)}.$$

T

he normalization constants μ_{v_j} , σ_{v_j} , $\max(v_j)$, $\min(v_j)$ have to be estimated from the *Training Set*. The same (training) constants are then applied during inference. Extracting constants from the *Test Set* is meaningless as the system has to infer with

feature post-processing

feature aggregation

feature aggregation:¹ compute *summary features* from feature series \Rightarrow **subfeatures**

■ reasons

- only one feature vector required per file
- data reduction
- characteristics of distribution or change over time contain additional info

■ examples

- *statistical descriptors*
 - ▶ mean, median, max, standard deviation
- *hand crafted*
 - ▶ anything that might be meaningful — periodicity, slope, ...

¹also compare *pooling* operation in machine learning

feature post-processing

feature aggregation

- could be for whole file or **texture window**:
split feature series in overlapping blocks of a few seconds length

- could be **hierarchical** process:
 - 1 compute subfeatures per window
 - 2 compute subfeatures of subfeature series
 - 3 (go to 1.)

feature post-processing

feature aggregation

- could be for whole file or **texture window**:
split feature series in overlapping blocks of a few seconds length
- could be **hierarchical** process:
 - 1 compute subfeatures per window
 - 2 compute subfeatures of subfeature series
 - 3 (go to 1.)

introduction

dimensionality reduction

■ problem

- many ML approaches cannot cope with large amounts of irrelevant features
- ML algorithms might degrade in performance

■ advantages

- reducing storage requirements
- reducing training complexity
- defying the “curse of dimensionality”

■ disadvantages

- additional workload for reduction
- adding an additional layer of model complexity

introduction

dimensionality reduction

■ problem

- many ML approaches cannot cope with large amounts of irrelevant features
- ML algorithms might degrade in performance

■ advantages

- reducing storage requirements
- reducing training complexity
- defying the “curse of dimensionality”

■ disadvantages

- additional workload for reduction
- adding an additional layer of model complexity

introduction

dimensionality reduction

■ problem

- many ML approaches cannot cope with large amounts of irrelevant features
- ML algorithms might degrade in performance

■ advantages

- reducing storage requirements
- reducing training complexity
- defying the “curse of dimensionality”

■ disadvantages

- additional workload for reduction
- adding an additional layer of model complexity

introduction

dimensionality issues

problems of high-dimensional data:

- increase in run-time
- overfitting
- curse of dimensionality
- required amount of training samples


introduction

dimensionality issues

problems of high-dimensional data:

- increase in run-time
- overfitting
- curse of dimensionality
- required amount of training samples

⇒ increasing number of input features may *decrease* classification performance



dimensionality issues

overfitting

■ overfitting:

- lack of training data
- overly complex model

⇒ model cannot be estimated properly


- required training set size depends on
 - ▶ classifier and its parametrization
 - ▶ number of classes
 - ▶ ...
- *rule of thumb:*
don't bother with training sets smaller than \mathcal{F}^2

dimensionality issues

overfitting

■ overfitting:

- lack of training data
 - overly complex model
- ⇒ model cannot be estimated properly



- required training set size depends on
 - ▶ classifier and its parametrization
 - ▶ number of classes

dimensionality issues

overfitting

■ overfitting:

- lack of training data
- overly complex model

⇒ model cannot be estimated properly

- required training set size depends on
 - ▶ classifier and its parametrization
 - ▶ number of classes
 - ▶ ...

- *rule of thumb:*

don't bother with training sets smaller than \mathcal{F}^2

dimensionality issues

curse of dimensionality

■ curse of dimensionality:

- increasing dimensionality leads to sparse training data
- neighborhoods of data points become less concentrated
- model tends to be harder to estimate in higher-dimensional space
- applies to distance-based algorithms

■ example (uniformly distributed data)

- identify region on axis covering **1% of data**
 - ▶ 1-D: 1% of x-axis
 - ▶ 2-D: 10% of x-axis/y-axis
 - ▶ 3-D: 21.5% of x-axis/y-axis/z-axis
 - ▶ 10-D: 63%
 - ▶ 100-D: 95%



dimensionality reduction

introduction

■ **feature subset selection:**

discard least helpful features

- high “discriminative” or descriptive power
- non-correlation to other features
- invariance to irrelevancies

■ feature space transformation:

map feature space

dimensionality reduction

introduction

■ feature subset selection:

discard least helpful features

- high “discriminative” or descriptive power
- non-correlation to other features
- invariance to irrelevancies


■ feature space transformation:

map feature space

feature subset selection

manual feature selection

example scatter
plots of pairs of
features in a
multi-class
scenario



feature subset selection

introduction

1 wrapper methods:

- *description*
 - ▶ use the “classifier” itself to evaluate feature performance
- *advantages*
 - ▶ taking into account feature dependencies
 - ▶ model dependency
- *disadvantages*
 - ▶ complexity
 - ▶ risk of overfitting

2 filter methods:

- *description*
 - ▶ use an objective function
- *advantages*
 - ▶ easily scalable
 - ▶ independent of classification algorithm
- *disadvantages*
 - ▶ no interaction with classifier
 - ▶ no feature dependencies

feature subset selection

introduction

1 wrapper methods:

- *description*
 - ▶ use the “classifier” itself to evaluate feature performance
- *advantages*
 - ▶ taking into account feature dependencies
 - ▶ model dependency
- *disadvantages*
 - ▶ complexity
 - ▶ risk of overfitting

2 filter methods:

- *description*
 - ▶ use an objective function
- *advantages*
 - ▶ easily scalable
 - ▶ independent of classification algorithm
- *disadvantages*
 - ▶ no interaction with classifier
 - ▶ no feature dependencies

feature subset selection

wrapper methods 1/2

1 single variable classification:

- *procedure*
 - ▶ evaluate each feature individually
 - ▶ choose the top N
- *complexity*
 - ▶ subsets to test: \mathcal{F}
- *challenges*
 - ▶ inter-feature correlation is not considered
 - ▶ feature combinations are not considered

2 brute force subset selection

- *procedure*
 - ▶ evaluate all possible feature combinations
 - ▶ choose the optimal combination
- *complexity*
 - ▶ subsets to test: $2^{\mathcal{F}}$

feature subset selection

wrapper methods 1/2

1 single variable classification:

- *procedure*
 - ▶ evaluate each feature individually
 - ▶ choose the top N
- *complexity*
 - ▶ subsets to test: \mathcal{F}
- *challenges*
 - ▶ inter-feature correlation is not considered
 - ▶ feature combinations are not considered

2 brute force subset selection

- *procedure*
 - ▶ evaluate all possible feature combinations
 - ▶ choose the optimal combination
- *complexity*
 - ▶ subsets to test: $2^{\mathcal{F}}$

feature subset selection

wrapper methods 2/2

4 sequential forward selection

- *procedure*

- ① init: empty feature subset $\mathcal{V}_s = \emptyset$
- ② find feature v_j maximizing objective function

$$v_j = \operatorname{argmax}_{\forall j | v_j \notin \mathcal{V}_s} J(\mathcal{V}_s \cup v_j)$$

- ③ add feature v_j to \mathcal{V}_s
- ④ go to step 2

5 sequential backward elimination

- *procedure*

- ① init: full feature set
- ② find feature v_j with the least impact on objective function
- ③ discard feature v_j
- ④ go to step 2

feature subset selection

wrapper methods 2/2

4 sequential forward selection

- *procedure*

- 1 init: empty feature subset $\mathcal{V}_s = \emptyset$
- 2 find feature v_j maximizing objective function

$$v_j = \underset{\forall j | v_j \notin \mathcal{V}_s}{\operatorname{argmax}} J(\mathcal{V}_s \cup v_j)$$

- 3 add feature v_j to \mathcal{V}_s
- 4 go to step 2

5 sequential backward elimination

- *procedure*

- 1 init: full feature set
- 2 find feature v_j with the least impact on objective function
- 3 discard feature v_j
- 4 go to step 2

feature space transformation

PCA introduction



■ objective

- map features to new coordinate system

$$\mathbf{u}(n) = \mathbf{T}^T \cdot \mathbf{v}(n)$$

- ▶ $\mathbf{u}(n)$: transformed features (same dimension as $\mathbf{v}(n)$)
- ▶ \mathbf{T} : transformation matrix ($\mathcal{F} \times \mathcal{F}$)

$$\mathbf{T} = \begin{bmatrix} \mathbf{c}_0 & \mathbf{c}_1 & \dots & \mathbf{c}_{\mathcal{F}-1} \end{bmatrix}$$

■ properties

- \mathbf{c}_0 points in the direction of highest *variance*
- variance concentrated in as few output components as possible
- \mathbf{c}_i orthogonal

$$\mathbf{c}_i^T \cdot \mathbf{c}_j = 0 \quad \forall i \neq j$$

- transformation is invertible

$$\mathbf{v}(n) = \mathbf{T} \cdot \mathbf{u}(n)$$

feature space transformation

PCA introduction

■ objective

- map features to new coordinate system

$$\mathbf{u}(n) = \mathbf{T}^T \cdot \mathbf{v}(n)$$

- ▶ $\mathbf{u}(n)$: transformed features (same dimension as $\mathbf{v}(n)$)
- ▶ \mathbf{T} : transformation matrix ($\mathcal{F} \times \mathcal{F}$)

$$\mathbf{T} = \begin{bmatrix} \mathbf{c}_0 & \mathbf{c}_1 & \dots & \mathbf{c}_{\mathcal{F}-1} \end{bmatrix}$$

■ properties

- \mathbf{c}_0 points in the direction of highest *variance*
- variance concentrated in as few output components as possible
- \mathbf{c}_i orthogonal



$$\mathbf{c}_i^T \cdot \mathbf{c}_j = 0 \quad \forall \quad i \neq j$$

- transformation is invertible

$$\mathbf{v}(n) = \mathbf{T} \cdot \mathbf{u}(n)$$

feature space transformation

PCA visualization





calculation of the transformation matrix

- 1 compute covariance matrix R

$$R = \mathcal{E}\{(V - \mathcal{E}\{V\})(V - \mathcal{E}\{V\})\}$$

feature space transformation

PCA visualization



calculation of the transformation matrix




- 1 compute covariance matrix R

$$R = \mathcal{E}\{(V - \mathcal{E}\{V\})(V - \mathcal{E}\{V\})\}$$

- 2 choose eigenvectors as axes for the new coordinate system

introduction

PCA example



introduction

PCA example

pca transformation matrix

$$\begin{bmatrix} -0.4187 & 0.3467 & -0.4569 & 0.4143 & -0.1271 & -0.5549 \\ -0.3908 & 0.1815 & 0.8136 & -0.0289 & 0.2060 & -0.3304 \\ -0.4516 & 0.3384 & 0.0859 & 0.2413 & -0.2919 & 0.7285 \\ -0.4337 & 0.1699 & -0.3337 & -0.7243 & 0.3747 & 0.0816 \\ 0.3802 & 0.5599 & -0.0381 & 0.2808 & 0.6622 & 0.1524 \\ 0.3679 & 0.6245 & 0.0956 & -0.4071 & -0.5267 & -0.1495 \end{bmatrix}$$

introduction

PCA example

introduction

PCA example

introduction

PCA example

pca transformation matrix

$$\begin{bmatrix} -0.4187 & 0.3467 & -0.4569 & 0.4143 & -0.1271 & -0.5549 \\ -0.3908 & 0.1815 & 0.8136 & -0.0289 & 0.2060 & -0.3304 \\ -0.4516 & 0.3384 & 0.0859 & 0.2413 & -0.2919 & 0.7285 \\ -0.4337 & 0.1699 & -0.3337 & -0.7243 & 0.3747 & 0.0816 \\ 0.3802 & 0.5599 & -0.0381 & 0.2808 & 0.6622 & 0.1524 \\ 0.3679 & 0.6245 & 0.0956 & -0.4071 & -0.5267 & -0.1495 \end{bmatrix}$$

summary

lecture content

■ feature matrix should be processed to adapt to task and classifier

- derive additional features
- aggregate features
- normalize features

■ derived features

- take existing features and “create” new ones

■ aggregate features: subfeatures

- combine blocks of features by computing, e.g., statistical features from them (mean, standard deviation, ...)
- subfeature vector is used as classifier input or as intermediate feature series

■ feature normalization

- avoid different value ranges might impacting classifier
- handle different feature distributions

