

Introduction to Audio Content Analysis

Module 5.0: Data, Data Splits, and Augmentation

alexander lerch

introduction

overview

corresponding textbook section

Section 5

■ lecture content

- data requirements
- data splits for train and test
- N-Fold cross-validation
- data augmentation

■ learning objectives

- understand the importance of data in machine learning
- define task-specific data requirements
- discuss possibilities of data augmentation
- implement N-Fold cross-validation in Python



introduction

overview

corresponding textbook section

Section 5

■ lecture content

- data requirements
- data splits for train and test
- N-Fold cross-validation
- data augmentation

■ learning objectives

- understand the importance of data in machine learning
- define task-specific data requirements
- discuss possibilities of data augmentation
- implement N-Fold cross-validation in Python



machine learning

data-driven

- derive classification parameters from data, e.g.,
⇒ learn feature distributions/separation metrics per class

- typical steps

- 1 define training set:** annotated results
- 2 normalize** training set
- 3 train** classifier
- 4 evaluate** classifier with test (or validation) set
- 5 (adjust** classifier settings, return to 4.)

machine learning

data-driven

- derive classification parameters from data, e.g.,
⇒ learn feature distributions/separation metrics per class

- typical steps

- 1 define training set:** annotated results
- 2 normalize** training set
- 3 train** classifier
- 4 evaluate** classifier with test (or validation) set
- 5 (adjust** classifier settings, return to 4.)

machine learning

data-driven

- derive classification parameters from data, e.g.,
⇒ learn feature distributions/separation metrics per class

- typical steps

- 1 **define training set:** annotated results
- 2 **normalize** training set
- 3 **train** classifier
- 4 **evaluate** classifier with test (or validation) set
- 5 (**adjust** classifier settings, return to 4.)

machine learning

data-driven

- derive classification parameters from data, e.g.,
⇒ learn feature distributions/separation metrics per class

- typical steps

- 1 **define training set:** annotated results
- 2 **normalize** training set
- 3 **train** classifier
- 4 **evaluate** classifier with test (or validation) set
- 5 (adjust classifier settings, return to 4.)

machine learning

data-driven

- derive classification parameters from data, e.g.,
⇒ learn feature distributions/separation metrics per class

- typical steps
 - 1 **define training set**: annotated results
 - 2 **normalize** training set
 - 3 **train** classifier
 - 4 **evaluate** classifier with test (or validation) set
 - 5 (**adjust** classifier settings, return to 4.)

data

requirements

what are important properties of our data





what are important properties of our data

■ representative

- represent all necessary factors of input data (e.g., range of genres, audio qualities, musical complexity, etc.)
- unbiased representation of class balance/label distribution

■ clean, non-noisy

- potential issues with subjective tasks

■ sufficient

- complex tasks/systems require lots of data

data

data split

- a bigger data set is commonly split in subsets
 - **training data** ($\approx 70 - 80\%$)
 - ▶ used to build the machine learning model
 - **validation data** ($\approx 10 - 15\%$)
 - ▶ used to tweak model parameters
 - **testing data** ($\approx 10 - 15\%$)
 - ▶ used to evaluate the model
 - ▶ needs to be **unseen!**

- no overlap between subsets!
 - also make sure that similar content (from one recording, album, artist, ...) is grouped into **one subset only**

data

N-Fold cross validation

- trying to utilize ALL data as both training and testing data
 - special case: Leave One Out CV
 - tends to be time-consuming
-
- 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

data

N-Fold cross validation

- trying to utilize ALL data as both training and testing data
 - special case: Leave One Out CV
 - tends to be time-consuming
-
- 1** split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

data

N-Fold cross validation

- trying to utilize ALL data as both training and testing data
 - special case: Leave One Out CV
 - tends to be time-consuming
-
- 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

data

N-Fold cross validation

- trying to utilize ALL data as both training and testing data
 - special case: Leave One Out CV
 - tends to be time-consuming
-
- 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

data

N-Fold cross validation

- trying to utilize ALL data as both training and testing data
 - special case: Leave One Out CV
 - tends to be time-consuming
-
- 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

data

N-Fold cross validation

- trying to utilize ALL data as both training and testing data
 - special case: Leave One Out CV
 - tends to be time-consuming
-
- 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

data

N-Fold cross validation

- trying to utilize ALL data as both training and testing data
 - special case: Leave One Out CV
 - tends to be time-consuming
-
- 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

classification

interaction of data, features, and classifier

■ training set

- training set too small, feature number too large
⇒ *overfitting*
- training set **too noisy**
⇒ *underfitting*
- training set **not representative**
⇒ *bad classification performance*

■ classifier

- classifier too complex
⇒ *overfitting*
- **poor classifier**
⇒ *bad classification performance*

■ features

- **poor features**
⇒ *bad classification performance*

classification

interaction of data, features, and classifier

■ training set

- training set too small, feature number too large
⇒ *overfitting*
- training set **too noisy**
⇒ *underfitting*
- training set **not representative**
⇒ *bad classification performance*

■ classifier

- classifier too complex
⇒ *overfitting*
- **poor classifier**
⇒ *bad classification performance*

■ features

- *poor features*
⇒ *bad classification performance*

classification

interaction of data, features, and classifier

■ training set

- training set too small, feature number too large
⇒ *overfitting*
- training set **too noisy**
⇒ *underfitting*
- training set **not representative**
⇒ *bad classification performance*

■ classifier

- classifier too complex
⇒ *overfitting*
- **poor classifier**
⇒ *bad classification performance*

■ features

- **poor features**
⇒ *bad classification performance*

data

augmentation

- if annotated data is insufficient, we can 'cheat' by increasing the amount of training data

⇒ **data augmentation**: apply irrelevant transforms to audio data

- *data segmentation*
 - ▶ treat audio snippets as separate observations
- *quality degradation*
 - ▶ add noise and distortion, limit bandwidth, etc.
- *audio effects*
 - ▶ apply reverb, etc.
- *changing pitch/tempo*
- *combine data*
 - ▶ mix different audio inputs together (if labels can be "mixed")
- *mask out parts of the signal*

data augmentation

- if annotated data is insufficient, we can 'cheat' by increasing the amount of training data

⇒ **data augmentation**: apply irrelevant transforms to audio data

- *data segmentation*
 - ▶ treat audio snippets as separate observations
- *quality degradation*
 - ▶ add noise and distortion, limit bandwidth, etc.
- *audio effects*
 - ▶ apply reverb, etc.
- *changing pitch/tempo*
- *combine data*
 - ▶ mix different audio inputs together (if labels can be "mixed")
- *mask out parts of the signal*

data augmentation

- if annotated data is insufficient, we can 'cheat' by increasing the amount of training data

⇒ **data augmentation**: apply irrelevant transforms to audio data

- *data segmentation*
 - ▶ treat audio snippets as separate observations
- *quality degradation*
 - ▶ add noise and distortion, limit bandwidth, etc.
- *audio effects*
 - ▶ apply reverb, etc.
- *changing pitch/tempo*
- *combine data*
 - ▶ mix different audio inputs together (if labels can be "mixed")
- *mask out parts of the signal*

data augmentation

- if annotated data is insufficient, we can 'cheat' by increasing the amount of training data

⇒ **data augmentation**: apply irrelevant transforms to audio data

- *data segmentation*
 - ▶ treat audio snippets as separate observations
- *quality degradation*
 - ▶ add noise and distortion, limit bandwidth, etc.
- *audio effects*
 - ▶ apply reverb, etc.
- *changing pitch/tempo*
- *combine data*
 - ▶ mix different audio inputs together (if labels can be "mixed")
- *mask out parts of the signal*

data augmentation

- if annotated data is insufficient, we can 'cheat' by increasing the amount of training data

⇒ **data augmentation**: apply irrelevant transforms to audio data

- *data segmentation*
 - ▶ treat audio snippets as separate observations
- *quality degradation*
 - ▶ add noise and distortion, limit bandwidth, etc.
- *audio effects*
 - ▶ apply reverb, etc.
- *changing pitch/tempo*
- *combine data*
 - ▶ mix different audio inputs together (if labels can be "mixed")
- *mask out parts of the signal*

data

augmentation

- if annotated data is insufficient, we can 'cheat' by increasing the amount of training data

⇒ **data augmentation**: apply irrelevant transforms to audio data

- *data segmentation*
 - ▶ treat audio snippets as separate observations
- *quality degradation*
 - ▶ add noise and distortion, limit bandwidth, etc.
- *audio effects*
 - ▶ apply reverb, etc.
- *changing pitch/tempo*
- *combine data*
 - ▶ mix different audio inputs together (if labels can be "mixed")
- *mask out parts of the signal*

data

augmentation

- if annotated data is insufficient, we can 'cheat' by increasing the amount of training data

⇒ **data augmentation**: apply irrelevant transforms to audio data

- *data segmentation*
 - ▶ treat audio snippets as separate observations
- *quality degradation*
 - ▶ add noise and distortion, limit bandwidth, etc.
- *audio effects*
 - ▶ apply reverb, etc.
- *changing pitch/tempo*
- *combine data*
 - ▶ mix different audio inputs together (if labels can be "mixed")
- *mask out parts of the signal*

summary

lecture content

■ data

- representative
- clean, non-noisy
- sufficient

■ data split

- train
- validation
- test

■ cross validation

- multiple runs with varying data splits
- maximum data utilization

